

```
In [2]: import dask.dataframe as dd
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [4]: df = dd.read_csv("kindle_reviews.csv", dtype={'reviewerName': 'object'})
```

```
In [6]: df.head()
```

```
Out[6]:   Unnamed: 0      asin  helpful  overall  reviewText  reviewTime  reviewerID  re
          0
0           0  B000F83SZQ  [0, 0]      5  I enjoy vintage books and movies so I enjoyed ...
1           1  B000F83SZQ  [2, 2]      4  This book is a reissue of an old one; the auth...
2           2  B000F83SZQ  [2, 2]      4  This was a fairly interesting read. It had ol...
3           3  B000F83SZQ  [1, 1]      5  I'd never read any of the Amy Brewster mysteri...
4           4  B000F83SZQ  [0, 1]      4  If you like period pieces - clothing, lingo, y...
```


```
In [8]: # Total reviews
total_reviews = df.shape[0].compute()
print(f"Total reviews: {total_reviews}")

# Unique products and reviewers
unique_products = df['asin'].nunique().compute()
unique_reviewers = df['reviewerID'].nunique().compute()
print(f"Unique products: {unique_products}")
print(f"Unique reviewers: {unique_reviewers}")
```

```
Total reviews: 982619  
Unique products: 61934  
Unique reviewers: 68223
```

```
In [9]: # Distribution of ratings  
rating_counts = df['overall'].value_counts().compute().sort_index()  
print("Rating counts:")  
print(rating_counts)  
  
# Plot rating distribution  
rating_counts.plot(kind='bar', color='skyblue')  
plt.title("Distribution of Product Ratings")  
plt.xlabel("Rating")  
plt.ylabel("Number of Reviews")  
plt.grid(True)  
plt.show()
```

```
Rating counts:  
overall  
1    23018  
2    34130  
3    96194  
4   254013  
5   575264  
Name: count, dtype: int64
```



```
In [16]: import ast  
  
# Function to safely parse the helpful column  
def parse_helpful_list(x):
```

```

try:
    values = ast.literal_eval(x)
    if isinstance(values, list) and len(values) == 2:
        return values[0], values[1]
except:
    return 0, 0
return 0, 0

# Map using Dask
helpful_parsed = df['helpful'].map(parse_helpful_list, meta=('helpful', 'object'))
helpful_df = pd.DataFrame(helpful_parsed.tolist(), columns=['helpful_num', 'helpful_den'])

# Reassign to original df
df = df.assign(helpful_num=helpful_df['helpful_num'], helpful_den=helpful_df['helpful_den'])

```

In [17]:

```
# Add helpfulness ratio safely
df['helpfulness_ratio'] = df.apply(
    lambda row: row.helpful_num / row.helpful_den if row.helpful_den > 0 else 0,
    axis=1,
    meta=('helpfulness_ratio', 'float')
)
```

In [20]:

```
# Mean helpfulness ratio by rating
helpfulness_by_rating = df.groupby('overall')['helpfulness_ratio'].mean().compute()
print(helpfulness_by_rating)


# Plot
helpfulness_by_rating.plot(kind='bar', color='lightgreen')
plt.title("Average Helpfulness Ratio by Rating")
plt.xlabel("Rating")
plt.ylabel("Avg Helpfulness Ratio")
plt.grid(True)
plt.show()
```

```

overall
1    0.369004
2    0.367690
3    0.360499
4    0.364891
5    0.369964
Name: helpfulness_ratio, dtype: float64

```

Average Helpfulness Ratio by Rating



```
In [21]: # Compute review text Length
df['review_length'] = df['reviewText'].str.len()

# Average review length by rating
avg_length_by_rating = df.groupby('overall')['review_length'].mean().compute()
print("Average review length by rating:")
print(avg_length_by_rating)


# Plot
avg_length_by_rating.plot(kind='bar', color='orange')
plt.title("Average Review Length by Rating")
plt.xlabel("Rating")
plt.ylabel("Average Length (characters)")
plt.grid(True)
plt.show()
```

Average review length by rating:

overall

| | |
|---|------------|
| 1 | 504.446564 |
| 2 | 611.172986 |
| 3 | 650.021436 |
| 4 | 704.150364 |
| 5 | 555.573188 |

Name: review_length, dtype: Float64



```
In [22]: # Top 10 most helpful reviews with at least 10 votes
top_helpful_reviews = df[df['helpful_den'] >= 10].nlargest(10, 'helpfulness_ratio')
print("Top 10 helpful reviews:")
top_helpful_reviews
```

Top 10 helpful reviews:

| | asin | reviewerID | overall | helpfulness_ratio |
|-------------|------------|----------------|---------|-------------------|
| 158 | B00JKZZ8K8 | AX0ZTNX6KVJ40 | 5 | 1.0 |
| 1274 | B00JLZ5J7E | A354MOJ56YP5Y8 | 5 | 1.0 |
| 1616 | B00JM2HDOS | A2I58O8E2AVBZM | 5 | 1.0 |
| 1910 | B00JMF2RPU | A2DRRMQPPWXB62 | 5 | 1.0 |
| 2070 | B00JMG2RU4 | A2UKWLNK4R52YN | 5 | 1.0 |
| 4358 | B00JO1PSX0 | A1XIQBGKvh6G8B | 5 | 1.0 |
| 6407 | B00JPJY43Q | A1YCJXNHARGRVR | 1 | 1.0 |
| 7025 | B00JPW10QC | AR35CG3TE4S4D | 3 | 1.0 |
| 7688 | B00QQD394 | A34HRQ4GESJ158 | 5 | 1.0 |
| 8473 | B00JRC0BCY | AAXG649HM1S5Q | 5 | 1.0 |

```
In [24]: # Group and aggregate
product_stats = df.groupby('asin').agg({'overall': ['mean', 'count']}).compute()
```

```

product_stats.columns = ['avg_rating', 'review_count']

# Filter and sort
top_products = product_stats[product_stats['review_count'] >= 50].sort_values(by='a
print("Top 10 best-rated products (min 50 reviews):")
top_products

```

Top 10 best-rated products (min 50 reviews):

Out[24]:

| | avg_rating | review_count |
|-------------------|------------|--------------|
| asin | | |
| B00ID7K5CA | 5.000000 | 70 |
| B00JM0RESG | 5.000000 | 59 |
| B00GN8K6CU | 4.975309 | 81 |
| B00DRN5X72 | 4.969388 | 98 |
| B0073999SU | 4.968254 | 63 |
| B00LDI41Z8 | 4.964912 | 57 |
| B00KP5P2EY | 4.963636 | 55 |
| B00IJBDBG2 | 4.961039 | 77 |
| B00DRERCRA | 4.957143 | 70 |
| B00C44PW5I | 4.951456 | 103 |

In [28]:

```

# Convert Unix time to datetime
df['review_datetime'] = dd.to_datetime(df['unixReviewTime'], unit='s')
df['review_year'] = df['review_datetime'].dt.year

# Reviews per year
reviews_per_year = df.groupby('review_year').size().compute().sort_index()

# Plot
reviews_per_year.plot(kind='line', marker='o', color='purple')
plt.title("Number of Reviews Over Years")
plt.xlabel("Year")
plt.ylabel("Number of Reviews")
plt.grid(True)
plt.show()

```

